

# ID2203 –Distributed Systems, Advanced

## Course Project – VT16 P3

### 1 Organisation

The course project consists of 4 parts. The first part is simply an introduction to Kompics and is optional if you have worked with Kompics before. The rest will be summarised in a final project report which is graded at the end of the course and forms the basis for the lab part of the course.

To motivate you to start working in time, there is an intermediate milestone after task **2.1**. You will have to submit a preliminary report in Canvas. Every group will be assigned another group to review their preliminary report and write a short review with feedback, which will also be submitted in Canvas.

The project is to be done in pairs, with clearly divided responsibilities. It is important to point out the contributions of each member in the final report.

Hand in both the report and the source code of your project in Canvas.

#### 1.1 Dates

**Tutorial Session** January 19th

**Preliminary Report** February 12th in Canvas

**Preliminary Report Peer Feedback Session** February 16th during the Exercise

**Code Working** March 3rd for the Project session

**Final Report** March 7th in Canvas

#### 1.2 Goals

The goal of the project is to implement and test a simple partitioned, distributed in-memory key-value store with linearisable operation semantics. You have significant freedoms in your choice of implementation, but you will need to motivate all your decisions in the report. Some of the issues to consider are:

- Networking Protocol

- Bootstrapping
- Group Membership
- Failure Detector
- Routing Protocol
- Replication Algorithm

You will also have to write verification scenarios for your implementation. Distributed algorithms are notoriously difficult to test and verify, so do not take this part of the tasks lightly.

*Note: A good labour distribution for the tasks might be to have one group member work on the algorithm implementations and the other one on the test scenarios after having agreed on the general architecture and interfaces.*

### 1.3 Requirements

For this lab you will need the following:

- Java SDK, version 7 or newer. Oracle Java is recommended but OpenJDK should work as well.
- Maven
- An IDE like Netbeans, Eclipse, IntelliJ is recommended, but any simple text-editor would be enough.

### 1.4 Time

Be sure to plan enough time for the project. The project will require a significant amount of code, and testing distributed systems is notoriously difficult and time intensive. It is recommended that you start as early as possibly to get a feel for Kompics and how long it takes to implement something in it.

## 2 Tasks

### 2.0 Introduction to Kompics (0 Points)

Implement all the *PingPong* examples from the Kompics tutorial at:  
<http://kompics.sics.se>.

This task is optional and does not give any points. However, if you haven't worked with Kompics before you should most definitely do it. If you have questions you can ask them during the tutorial exercise session.

## 2.1 Infrastructure (20 Points)

For this task you have to design the basic infrastructural layers for the key-value store. Your system should support a partitioned key-space of some kind (e.g. hash-partitioned strings or range-partitioned integers). The partitions need to be distributed over the available nodes such that all value are replicated with a specific replication degree  $\delta$ . You are free to keep  $\delta$  as a configuration value or hardcode it, as long as it fulfils the requirements of your chosen replication algorithm (task 2.2), so plan ahead.

For this task you need to be able to set up the system, assign nodes to partitions and replication groups, lookup (preloaded) values<sup>1</sup> from an external system (client), and detect (but not handle) failures of nodes. Additionally, you should be able to broadcast information within replication groups and possibly across.

On the verification side you have to write simulator scenarios that test all the features mentioned above. You should also be able to run in a simple deployment (it's fine to run multiple JVMs on the same machine and consider them separate nodes).

For the report describe and motivate all your decisions and tests.

## 2.2 KV-Store (20 Points)

After you have gotten the basic infrastructure working, you have to add a *PUT*(*key*, *value*) operation, that updates (or adds) a value at a key, to the *GET* from the previous task. As mentioned in the goals, the operations should fulfil the linearisable property, so make sure choose you the right replication algorithm.

For more points, also implement a compare-and-swap (*CAS*(*key*, *referenceValue*, *newValue*) : *success?*) operation that compares the current value at the key to a given reference value and only updates with the new value if the old value and the reference value are the same.

As before, be sure to write test scenarios for the simulator that check the correctness of your implementation. Especially be very careful to explain how you are verifying the linearisability of the store.

## 2.3 Reconfiguration (10 Bonus Points)

At this point your store is fairly static and can't really deal with node failures apart from complaining about them. For this task you should implement reconfiguration support for your replication groups and your routing protocol. You should be able to deal with both nodes leaving the system (treat a voluntary leave the same way as a failure, for simplicity) and new nodes joining the system. There are many ways to interpret the semantics of this, including making some of the partitions unavailable while they are under-replicated. Any approach you take is acceptable as long as you document it properly in the report. However, you have to make sure that reconfiguration does not violate linearisability for *correct* nodes.

---

<sup>1</sup>This is effectively a *GET*(*key*) : *value* operation.

To get full points for this task you'll have to write test scenarios that reconfigure the system and verify that for all correct nodes the operations are still linearisable.